

Index

What is an Index?

An index in a database is a data structure that improves the speed of data retrieval operations on a table at the cost of additional writes and storage space to maintain it. Indexes are used to quickly locate data without having to search every row in a database table every time a database table is accessed.

Types of Indexes

1. **Primary Index:**

- Automatically created when a primary key is defined.
- Ensures that the values in the key are unique and not null.

2. **Unique Index:**

- Ensures that all the values in the indexed column are unique.
- Created using the UNIQUE keyword.

3. **Non-Unique Index:**

- Allows duplicate values in the indexed column.
- Created without the UNIQUE keyword.

4. **Composite Index:**

- An index on multiple columns of a table.
- Useful for queries involving multiple columns in the WHERE clause.

5. **Full-Text Index:**

- Special type of index used for full-text searches.
- Suitable for searching large text fields.

6. **Spatial Index:**

- Used for spatial data types.
- Enhances performance for spatial queries.

How Indexes Work

Indexes are typically implemented using B-trees or B+ trees. Here's how they work in general:

- **B-tree Structure:** The B-tree structure keeps data sorted and allows searches, sequential access, insertions, and deletions in logarithmic time.
- **B+ Tree:** A variation of the B-tree, used in databases to improve the efficiency of range queries.

When a query is executed, the database engine uses the index to quickly locate the data, reducing the number of rows that need to be scanned.

Benefits of Using Indexes

1. Improved Query Performance:

- Speeds up the retrieval of rows by reducing the number of disk I/O operations.
- Efficient for large tables.

2. Faster Sorting:

- Helps in quickly sorting rows using the indexed columns.

3. Quick Access to Rows:

- Enhances performance for SELECT, JOIN, and WHERE clauses.

Costs of Using Indexes

1. Increased Storage:

- Indexes consume additional disk space.

2. Slower Write Operations:

- INSERT, UPDATE, and DELETE operations can be slower because the index must be updated whenever the data changes.

3. **Maintenance Overhead:**

- Requires ongoing maintenance, especially for frequently updated tables.

When to Use Indexes

1. **Frequent Searches:**

- Use indexes on columns that are frequently searched.

2. **Sorting:**

- Useful for columns involved in ORDER BY clauses.

3. **Joins:**

- Indexes on foreign keys can speed up join operations.

4. **Unique Constraints:**

- Enforce uniqueness on columns.

When Not to Use Indexes

1. **Small Tables:**

- For small tables, full table scans are usually more efficient than using indexes.

2. **Frequent Updates:**

- Tables with frequent insertions, updates, or deletions may suffer from the overhead of maintaining indexes.

3. **Columns with High Cardinality:**

- Columns with a high number of unique values might not benefit much from indexing.

Step 1: Create the Table

First, let's create a table named Employees.

```
CREATE TABLE Employees (  
    EmployeeID INT PRIMARY KEY AUTO_INCREMENT,  
    FirstName VARCHAR(50),  
    LastName VARCHAR(50),  
    Position VARCHAR(50),  
    Salary DECIMAL(10, 2),  
    HireDate DATE  
);
```

Step 2: Insert Records

Next, let's insert some records into the Employees table.

```
INSERT INTO Employees (FirstName, LastName, Position, Salary,  
HireDate) VALUES  
( 'Rahul', 'Sharma', 'Manager', 60000, '2020-01-15'),  
( 'Anjali', 'Verma', 'Developer', 50000, '2019-02-10'),  
( 'Raj', 'Kumar', 'Designer', 45000, '2021-03-22'),  
( 'Priya', 'Patel', 'Tester', 40000, '2018-07-30');
```

Step 3: Create Indexes

Indexes are used to speed up the retrieval of data from a database table. Here are examples of creating different types of indexes:

Creating an Index on a Single Column

```
CREATE INDEX idx_lastname ON Employees (LastName);
```

Creating a Composite Index on Multiple Columns

```
CREATE INDEX idx_lastname_position ON Employees (LastName, Position);
```

Complete SQL Script

Here is the complete SQL script combining all the steps:

-- Step 1: Create the Employees table

```
CREATE TABLE Employees (  
    EmployeeID INT PRIMARY KEY AUTO_INCREMENT,  
    FirstName VARCHAR(50),  
    LastName VARCHAR(50),  
    Position VARCHAR(50),  
    Salary DECIMAL(10, 2),  
    HireDate DATE  
);
```

-- Step 2: Insert records into the Employees table

```
INSERT INTO Employees (FirstName, LastName, Position, Salary,  
HireDate) VALUES  
( 'Rahul', 'Sharma', 'Manager', 60000, '2020-01-15'),  
( 'Anjali', 'Verma', 'Developer', 50000, '2019-02-10'),  
( 'Raj', 'Kumar', 'Designer', 45000, '2021-03-22'),  
( 'Priya', 'Patel', 'Tester', 40000, '2018-07-30');
```

-- Step 3: Create indexes

-- Create an index on the LastName column

```
CREATE INDEX idx_lastname ON Employees (LastName);
```

-- Create a composite index on the LastName and Position columns

```
CREATE INDEX idx_lastname_position ON Employees (LastName,  
Position);
```

Verifying the Indexes

To verify that the indexes have been created, you can use the following SQL statement:

SHOW INDEX FROM Employees;

This will show a list of all indexes created on the Employees table, including their types and the columns they cover.

-- Drop the index on the LastName column

DROP INDEX idx_lastname ON Employees;

-- Drop the composite index on the LastName and Position columns

DROP INDEX idx_lastname_position ON Employees;